

## Document

### Chaining return-into-libc calls

1. “esp lifting” method
2. Frame faking
3. Inserting null bytes

### PaX features

1. PaX basics
2. PaX and return-into-libc exploits
3. PaX and mmap base randomization

### The dynamic linker's dl-resolve() function

1. A few ELF data types
2. A few ELF data structures
3. How dl-resolve() is called from PLT

### Defeating PaX

1. Requirements
2. Building the exploit

### Misc

1. Portability
2. Other types of vulnerabilities
3. Other non-exec solutions
4. Improving existing non-exec schemes

## 5. The versions used

## Practice

Babystack from Octl 2018

PREPARE:

- JMPREL
- SYMTAB
- STRTAB

```
STRTAB      0x80484f8 the location of string table (type char *)
SYMTAB      0x8048268 the location of symbol table (type Elf32_Sym*)
```

```
typedef uint32_t Elf32_Addr ;
typedef uint32_t Elf32_Word ;
typedef struct
{
    Elf32_Addr r_offset ; /* Address */
    Elf32_Word r_info ; /* Relocation type and symbol index */
} Elf32_Rel ;
#define ELF32_R_SYM(val) ((val) >> 8)
#define ELF32_R_TYPE(val) ((val) & 0xff)
```

```
typedef struct
{
    Elf32_Word st_name ; /* Symbol name (string tbl index) */
    Elf32_Addr st_value ; /* Symbol value */
    Elf32_Word st_size ; /* Symbol size */
    unsigned char st_info ; /* Symbol type and binding */
    unsigned char st_other ; /* Symbol visibility under glibc>=2.2 */
    Elf32_Section st_shndx ; /* Section index */
} Elf32_Sym ;
```

```
// call of unresolved read(0, buf, 0x100)
_dl_runtime_resolve(link_map, rel_offset) {
    Elf32_Rel * rel_entry = JMPREL + rel_offset ;
    Elf32_Sym * sym_entry = &SYMTAB [ ELF32_R_SYM ( rel_entry -> r_info )];
    /* Check... */
    char * sym_name = STRTAB + sym_entry -> st_name ;
    _search_for_symbol(link_map, sym_name);
    // invoke initial read call now that symbol is resolved
    read(0, buf, 0x100);
}
```

The following is the simplified dl-resolve() algorithm:

- 1) calculate some\_func's relocation entry  
Elf32\_Rel \* reloc = JMPREL + reloc\_offset;
- 2) calculate some\_func's symtab entry  
Elf32\_Sym \* sym = &SYMTAB[ ELF32\_R\_SYM (reloc->r\_info) ];
- 3) sanity check  
assert (ELF32\_R\_TYPE(reloc->r\_info) == R\_386\_JMP\_SLOT);
- 4) late glibc 2.1.x (2.1.92 for sure) or newer, including 2.2.x, performs another check. if sym->st\_other & 3 != 0, the symbol is presumed to have been resolved before, and the algorithm goes another way (and probably ends with SIGSEGV in our case). We must ensure that sym->st\_other & 3 == 0.
- 5) if symbol versioning is enabled (usually is), determine the version table  
index  
uint16\_t ndx = VERSYM[ ELF32\_R\_SYM (reloc->r\_info) ];  
and find version information  
const struct r\_found\_version \*version =&l->l\_versions[ndx];  
where l is the link\_map parameter. The important part here is that ndx must be a legal value, preferably 0, which means "local symbol".
- 6) the function name (an asciiz string) is determined:  
name = STRTAB + sym->st\_name;
- 7) The gathered information is sufficient to determine some\_func's address. The results are cached in two variables of type Elf32\_Addr, located at reloc->r\_offset and sym->st\_value.
- 8) The stack pointer is adjusted, some\_func is called.

Note: in case of glibc, this algorithm is performed by the fixup() function, called by dl-runtime-resolve().

For demonstration purposes only, let us suppose that:

- JMPREL @ 0x0
- SYMTAB @ 0x100
- STRTAB @ 0x200
- controllable area @ 0x300

We need to craft our Elf32\_Rel and Elf32\_Sym somewhere within the controllable area and provide a rel\_offset such that the resolver reads our special forged structures. Let's suppose that the controllable (stack after pivotation ??? ) are has the following layout.

r_offset	+-----+	
	GOT	0x300
r_info	0x2100	0x304
alignment	AAAAAAA	0x308
st_name	0x120	0x310
st_value	0x0	
st_size	0x0	
others	0x12	
sym_string	"syst	0x320
	em\x00"	
	+-----+	

When `_dl_runtime_resolve ( link_map , 0x300)` is called, the 0x300 offset is used to get the `Elf32_Rel* rel = JMPREL + 0x300 == 0x300`.

Secondly, the `Elf32_Sym` is accessed using the `r_info` field from 0x304.

`Elf32_Sym* sym = &SYMTAB[(0x2100 >> 8)] == 0x310`.

The last step is to compute the address of the symbol string. This is done by

adding `st_name` to `STRTAB` : `const char *name = STRTAB + 0x120 == 0x320`.

Note that `SYMTAB` access its entries as an array, therefore `ELF32_sym` should be aligned to 0x10 bytes. Now that we control `st_name`, we can basically force the resolver to relocate `system` and call `system('sh')` to own the system :)

Writing the payload should be easy now that we have a clear image of the forged memory layout.

## Ret-into-dl\_resol x64

What is the difference between the x86 and x64 Elf32\_Rel, Elf32\_Sym structure instead of using a Elf64\_Sym structure Elf64\_Rela,.

The important thing here is that the change in the size of the structure.

- The size of the structure (8 byte) → Elf32\_Rel Elf64\_Rela structures, size (24-byte)
- The size of the Elf32\_Sym structure (16 byte) → Elf64\_Sym The size of the structure (24-byte)

Because of this reloc\_offset value is an array of non-offset address Elf64\_Rela structures should be an index.

```
typedef uint32_t Elf64_Word;
typedef uint64_t Elf64_Xword;
typedef int64_t Elf64_Sxword;
typedef uint64_t Elf64_Addr;
typedef uint16_t Elf64_Half;

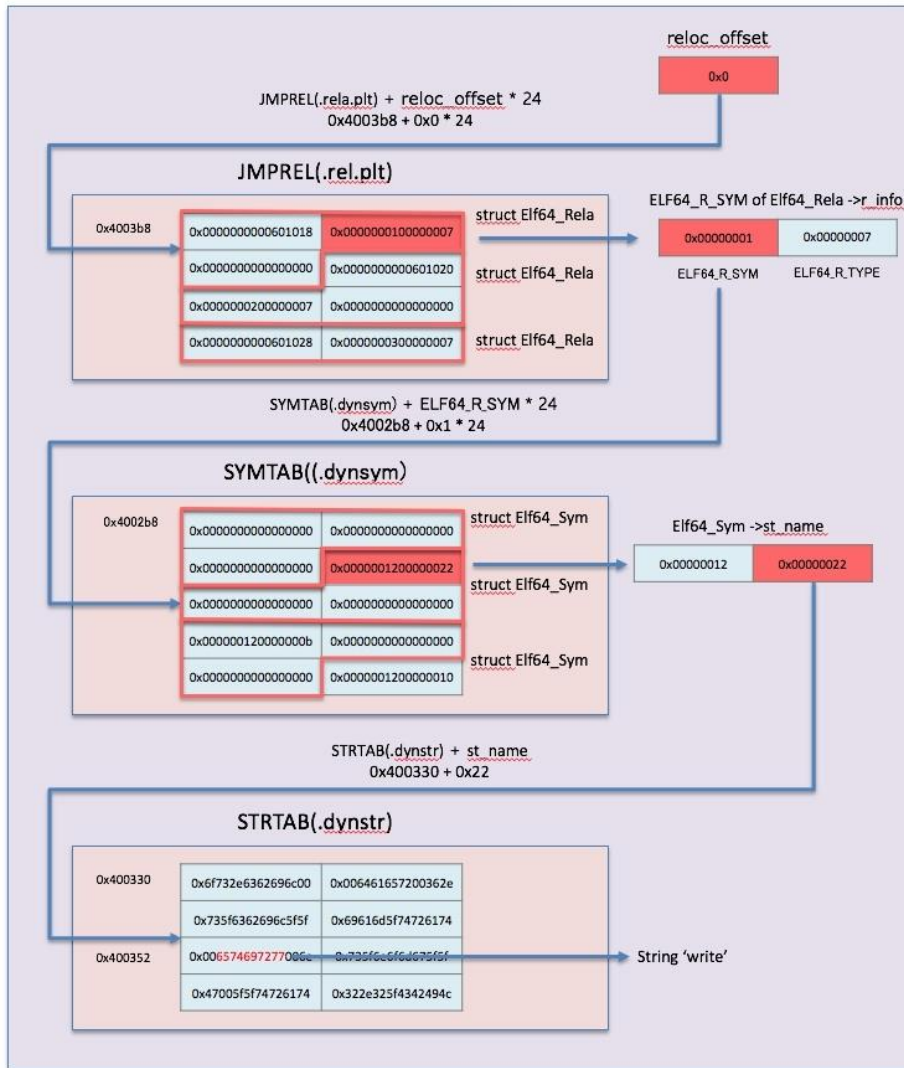
typedef struct
{
    Elf64_Addr    r_offset;           /* Address */
    Elf64_Xword   r_info;             /* Relocation type and symbol index */
} Elf64_Rel;

typedef struct
{
    Elf64_Sxword  r_addend;           /* Addend */
} Elf64_Rela;

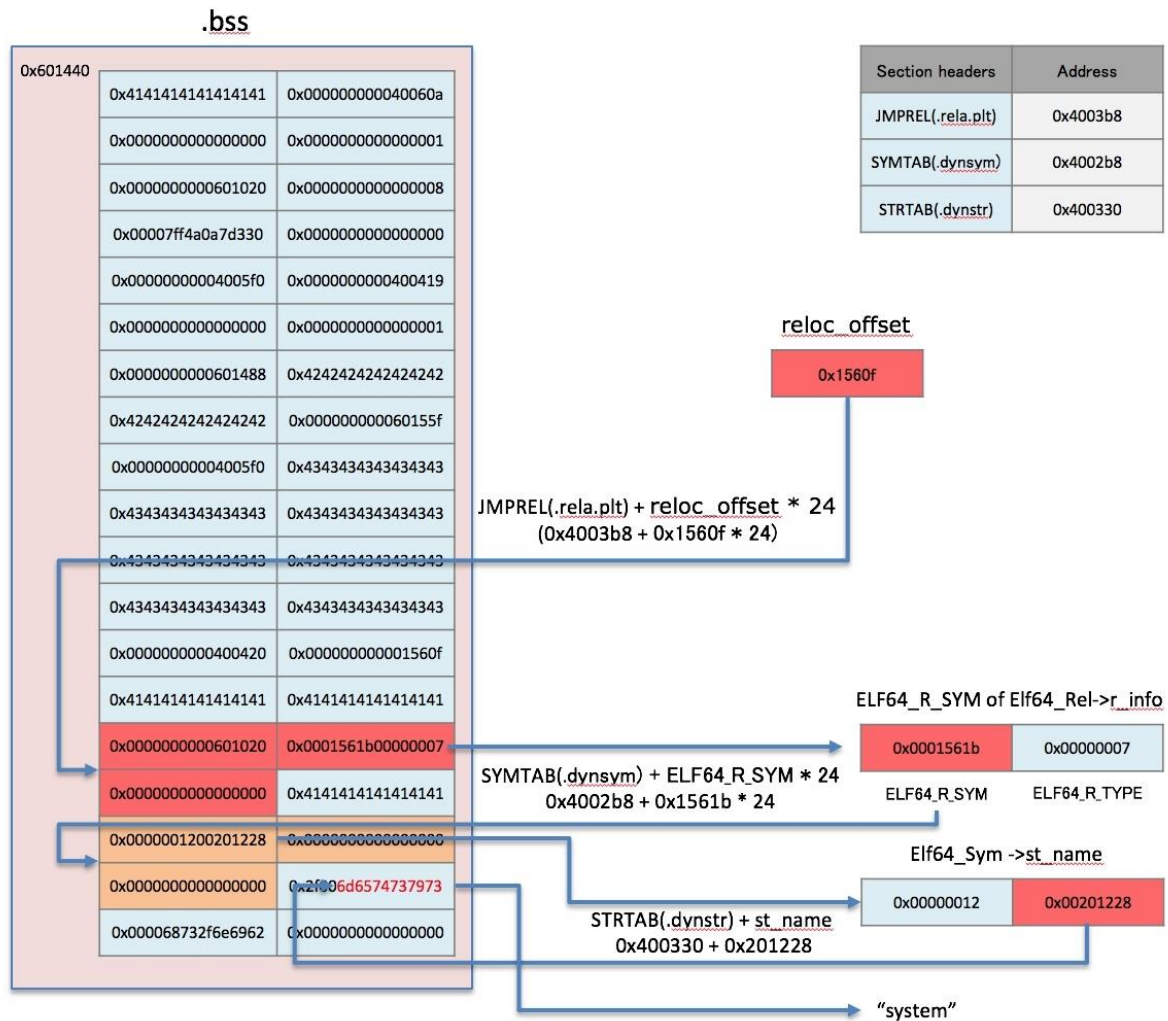
typedef struct
{
    Elf64_Word     st_name;           /* Symbol name (string tbl index) */
    unsigned char  st_info;           /* Symbol type and binding */
    unsigned char  st_other;          /* Symbol visibility */
    Elf64_Half     st_shndx;          /* Section index */
    Elf64_Addr     st_value;          /* Symbol value */
    Elf64_Xword    st_size;           /* Symbol size */
} Elf64_Sym;
```



## \_dl\_fixup()



Section headers	Address
JMPREL(.rela.plt)	0x4003b8
SYMTAB(.dynsym)	0x4002b8
STRTAB(.dynstr)	0x400330



## Return-to-csu

### Command

```
objdump -M intel -d ./prog
```

```

4005f0: 4c 89 ea      mov     rdx,r13
4005f3: 4c 89 f6      mov     rsi,r14
4005f6: 44 89 ff      mov     edi,r15d
4005f9: 41 ff 14 dc   call    QWORD PTR [r12+rbx*8]
4005fd: 48 83 c3 01   add     rbx,0x1
400601: 48 39 eb      cmp     rbx,rbp
400604: 75 ea        jne     4005f0 <__libc_csu_init+0x40>
400606: 48 83 c4 08   add     rsp,0x8
40060a: 5b          pop     rbx
40060b: 5d          pop     rbp
40060c: 41 5c        pop     r12
40060e: 41 5d        pop     r13
400610: 41 5e        pop     r14
400612: 41 5f        pop     r15
400614: c3          ret

```



We can take advantage of the gadget of `__libc_csu_init`

- We can control rdx, rsi, edi with r13, r14, r15 register

The important thing is that in 64bit, the argument is saved in register.

Operand size (bits)	Argument number					
	1	2	3	4	5	6
64	%rdi	%rsi	%rdx	%rcx	%r8	%r9
32	%edi	%esi	%edx	%ecx	%r8d	%r9d
16	%di	%si	%dx	%cx	%r8w	%r9w
8	%dil	%sil	%dl	%cl	%r8b	%r9b

**Figure 3.28** Registers for passing function arguments. The registers are used in a specified order and named according to the argument sizes.

63	31	15	7	0	
%rax	%eax	%ax	%al		Return value
%rbx	%ebx	%bx	%bl		Callee saved
%rcx	%ecx	%cx	%cl		4th argument
%rdx	%edx	%dx	%dl		3rd argument
%rsi	%esi	%si	%sil		2nd argument
%rdi	%edi	%di	%dil		1st argument
%rbp	%ebp	%bp	%bpl		Callee saved
%rsp	%esp	%sp	%spl		Stack pointer
%r8	%r8d	%r8w	%r8b		5th argument
%r9	%r9d	%r9w	%r9b		6th argument
%r10	%r10d	%r10w	%r10b		Caller saved
%r11	%r11d	%r11w	%r11b		Caller saved
%r12	%r12d	%r12w	%r12b		Callee saved
%r13	%r13d	%r13w	%r13b		Callee saved
%r14	%r14d	%r14w	%r14b		Callee saved
%r15	%r15d	%r15w	%r15b		Callee saved

#### Reference:

1. <http://lia.deis.unibo.it/Courses/SicurezzaM1011/BufferOverflow.pdf>
2. Chapter 7: Linker CSAPP
3. <https://gist.github.com/ricardo2197/8c7f6f5b8950ed6771c1cd3a116f7e62>
4. <http://phrack.org/issues/58/4.html?fbclid=IwAR3H6sjm3ouNN6rPIUuo2n1Rlu9JB-00SggAIP2fzokcGtT-hYOe0fj6jsQ>
5. <https://kileak.github.io/ctf/2018/Octf-qual-babystack/>
6. <http://inaz2.hatenablog.com/entry/2014/07/15/023406>
7. <http://inaz2.hatenablog.com/entry/2014/07/27/205322>
8. <https://www.lazenca.net/pages/viewpage.action?pageId=19300744>